

Wire Cell Event Reconstruction Software for LArTPC Detectors

Brett Viren

for the BNL Wire Cell Group

Physics Department



BNL CSC Seminar 2015 Sep 29

Outline

Introduction

Technique

- Data Preparation

- Imaging Of Activity

- Pattern Recognition

- Physics

Software Design

- Bee Display

- Prototype

- Toolkit

Future Plans

- Algorithm Development

- Bee 2.0

- Parallel Wire Cell

Who, What, How, Why

who The **Wire Cell** development is centered in the **Electronic Detector Group** of the **BNL Physics Department**.

what Experimental study of the **properties of neutrinos**,

how using **Liquid Argon Time Projection Chambers** (LArTPC) and **accelerator-neutrino beams**

why in order to reach these **Physics objectives**:

- Discover and measure **CP symmetry violation** in neutrino interactions.
- Determine **neutrino mass ordering**.
- Precisely measure **neutrino oscillation parameters**.
- Potentially observe **neutrino bursts from supernova**.
- Search for **proton decay** and **sterile neutrinos**.

Why Liquid Argon TPC

These goals require neutrino detectors to be:

- big** active mass of 10 – 100 kt's (mitigate tiny neutrino cross sections).
- efficient** to capture neutrinos and separate signal from background.
- accurate** and precise to measure energy, neutrino flavor and event type.
- quiet** low background environment shielded against cosmic- μ and natural radioactivity, \Rightarrow deep underground.

Two main, excellent technologies:

Water Cherenkov big++, **efficient+**, accurate++, quiet++
simple, mature technology, modest resolving power.


LArTPC **big+**, efficient++, accurate++, quiet++,
high-resolution, R&D still improving, excellent future potential.

Many challenges still to overcome, but **LArTPC's** promises make it the **chosen technology for neutrino experiments in the US.**

LArTPC Experiments - ICARUS


The origin of LArTPC technology for Neutrinos: C. Rubbia, 1977 led to **ICARUS**, the first, large-scale LArTPC.

- 2×300 t modules.
- Took data at Gran Sasso tunnel, Italy from CERN neutrino beam.
- Moving to Fermilab as part of the **Short-Baseline Neutrino** Program.



LArTPC Experiments - MicroBooNE

Just started taking data at Fermilab!




- 85 t fiducial mass.
- 8256 channels
- 3 mm wire pitch.
- Investigate *low energy excess puzzle*, look for sterile- ν , measure ν -Ar cross sections.

MicroBooNE is the initial test bed for Wire Cell reconstruction.

LArTPC Experiments - DEEP UNDERGROUND NEUTRINO EXPERIMENT

“International **mega-science** project”



Three stages of LArTPC detectors:


- ① “**35ton**” prototype (at FNAL) commissioning now, running early 2016.
- ② Full-scale “**protoDUNE**” (at CERN) 2017 with π , K , p beam tests.
- ③ Full “**DUNE**” far detector underground in South Dakota ~2025.
 - **40 kt fiducial mass** in 4 modules.
 - 5 mm wire pitch, **1.5M channels**.

DUNE design “features” driven by **detector size** and **underground location** \Rightarrow some extra reconstruction challenges (two sided anode planes + wrapped readout wires = extra ambiguities).

Next: primer on how LArTPC works.


Illustration by Bo Yu (BNL)

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




Charged particles traverse the detector at relativistic speeds. This leaves behind **ionized argon** and **electrons** in their tracks.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




Ion and electron pairs drift off in opposite directions in the electric field applied between **anode wires** and the cathode plane.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




Here, we watch the electrons drift. They drift with a speed of $1.6 \text{ mm}/\mu\text{s}$ (for nominal E-field of 500 V/cm).

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




As the electrons approach the wires they induce signals. Wires in each of **three planes provide independent measure** of the **drifting charge**.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




There are **two** planes of **induction wires** (waveform from “V” plane in blue). Wires see a **bipolar signal** as the charge first drifts toward and then away.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




The final, **third** plane **collects the drifting charge** and thus sees a **unipolar signal** (in **red waveform**).

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




The exact **signal response** is complex and depends on details of the local electric field and readout electronics.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




Charge waveform from each wire is digitized as a function of time. Typical digitization “tick” is $0.5\mu\text{s}$, wire spacing is 3 to 5 mm.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




Digitized waveforms are **deconvolved** with a **response function** and a **noise filter** to infer the amount of drifted charge near the wires at any given time.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




As a batch of electrons drift, it **diffuses** in both the **longitudi-nal** and **transverse** directions. LAr impurities also **attenuate** the charge.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




At the maximum drift distance allowed by the detector, diffusion can be as much as $1\ \mu\text{s}$ longitudinally and $2.5\ \text{mm}$ transverse.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber



Still drifting...

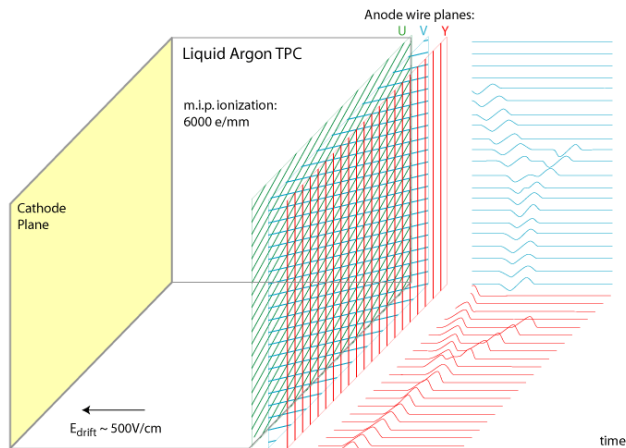
LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber



Still drifting... It's a slow detector!


$1.6 \text{ mm}/\mu\text{s}$ drift speed $\times \sim \text{meter drifts} \Rightarrow \text{few ms to readout.}$

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




Still drifting...

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




Finally, each wire plane reads out one 2D view (time vs. wire-pitch direction) of the original pattern of ionizing particles.
The drifting charge has three independent measures.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber



Combining three 2D views can give a tomographic reconstruction of the location and amount of energy deposition of the original particle trajectories.

LArTPC = Liquid **A**rgon **T**ime **P**roportional **C**hamber




Traditional methods work in 2D first, then combine views by assuming/asserting a particular event topology.

Wire Cell goes directly to 3D imaging with no such assumptions.

Some Data Numerology

LArTPC can produce **prodigious quantities** of **high-resolution** data from **huge detector volumes**:

- $10^4 - 10^6$ channels
- 2MHz @ 12 bit FADC digitization
- each “event” spans several milliseconds



0.5 μ s waveform digitization.

Two general DAQ readout strategies:

Full Stream reads out entire waveform (**MicroBooNE**)

- **30GB/s in 120 MB “events”.**
- DUNE would produce 5 TB/s in 25 GB “events”!

Zero Supression drops waveform chunks below a threshold (**DUNE**)

- Threshold based on noise ($E_{equiv} < 0.5 \text{ MeV/wire}$)
- 2.5 MB/event \rightarrow **100's TB/year**
- caveat: must veto ^{39}Ar decay in DAQ or accept **50 PB/year**

Introduction

Technique

Data Preparation

Imaging Of Activity

Pattern Recognition

Physics

Software Design

Future Plans

Wire Cell Reconstruction Method

Four main parts:

1 Data Preparation

- Construct wire geometry and **associated cells**.
- Read framed data stream, form **time slices**

2 Imaging of Activity


- The heart of the Wire Cell technique
- Identify the likely “hit” **cells** in the time slice.

3 Pattern Recognition

- Application of Wire Cell imaging
- Cluster** activity in space and across time slices.
- Categorize**: track, shower, etc.

4 Physics Quantities

- Determine particle ID and kinematics of tracks/showers.



Wire Cell Reconstruction Method

Four main parts:

1 Data Preparation

- Construct wire geometry and **associated cells**.
- Read framed data stream, form **time slices**

2 Imaging of Activity


- The heart of the Wire Cell technique
- Identify the likely “hit” **cells** in the time slice.

3 Pattern Recognition

- Application of Wire Cell imaging
- Cluster** activity in space and across time slices.
- Categorize**: track, shower, etc.

4 Physics Quantities

- Determine particle ID and kinematics of tracks/showers.



Wire Cell Reconstruction Method

Four main parts:

1 Data Preparation

- Construct wire geometry and **associated cells**.
- Read framed data stream, form **time slices**

2 Imaging of Activity


- The heart of the Wire Cell technique
- Identify the likely “hit” **cells** in the time slice.

3 Pattern Recognition

- Application of Wire Cell imaging
- Cluster** activity in space and across time slices.
- Categorize**: track, shower, etc.

4 Physics Quantities

- Determine particle ID and kinematics of tracks/showers.



Wire Cell Reconstruction Method

Four main parts:

1 Data Preparation

- Construct wire geometry and **associated cells**.
- Read framed data stream, form **time slices**

2 Imaging of Activity


- The heart of the Wire Cell technique
- Identify the likely “hit” **cells** in the time slice.

3 Pattern Recognition

- Application of Wire Cell imaging
- **Cluster** activity in space and across time slices.
- **Categorize**: track, shower, etc.

4 Physics Quantities

- Determine particle ID and kinematics of tracks/showers.



Wire Cell Reconstruction Method

Four main parts:

1 Data Preparation

- Construct wire geometry and **associated cells**.
- Read framed data stream, form **time slices**

2 Imaging of Activity


- The heart of the Wire Cell technique
- Identify the likely “hit” **cells** in the time slice.

3 Pattern Recognition

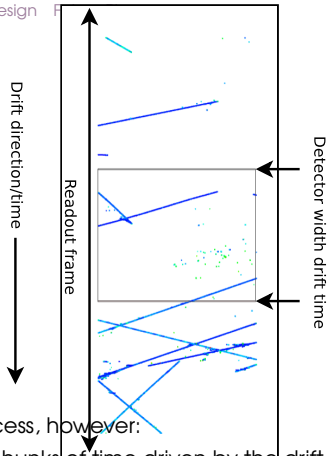
- Application of Wire Cell imaging
- Cluster** activity in space and across time slices.
- Categorize**: track, shower, etc.

4 Physics Quantities

- Determine particle ID and kinematics of tracks/showers.



Framing





Wire Cell is a streaming process, however:

- DAQ writes out large chunks of time driven by the drift time across the detector,
- Seamlessly capture any early/late activity encroaching around the trigger time.
- Triggerless running, to capture any low threshold early/late activity.

To Wire Cell, frames are “just an artifact” of the DAQ, wire cell wants time slices →


Time Slicing



- Slice duration is chosen to **match electronics shaping time**: 4 FADC "ticks" = $2\ \mu\text{s}$.
- Shows "true" (simulated) hits **red circles**.
- Select **wires hit** in the slice.
→ hit = FADC charge-in-slice above a threshold.
- Potential "**cells**" holding charge.

Tiling

Zoom in on the wires and their associated (constructed) cells.



MicroBooNE geometry, grey wires, colored cells.

Cell construction:

- Define small, 2D regions near *approximate triple crossings* of one wire from each of the three planes.
- A “cell” is the association of this region with the three wires.
- Cells completely tile the plane (time slice): no gaps, no overlaps.

Awkward consequences of current construction:

- Two cell types: *wire-centered* and *gap-centered* cells.
- In general (eg, DUNE), cell shapes and sizes not uniform nor regular, depend on wire plane pitches, angles and offsets.

The heart of the Wire Cell concept: if all three triple-crossing **wires** are “hit”, the associated **cell** likely contains drifted charge.

Good wire v3 measures no charge,
 \therefore all its cells must not be hit.

Ambiguous multiple cells measured by same wire.
How much charge is in **c6**???


In some cases ambiguity can not be resolved at the cell level!

Solution Attempt 1 - cell-level

The charge expected on the **wires** (\vec{w}) can be calculated given (the unknown) charge in the **cells** (\vec{c}):

$$\vec{w} = \mathbf{G}_{wc} \vec{c}$$

\mathbf{G}_{wc} is the **wire-cell adjacency matrix**, purely geometrical and perfectly known, function of detector design.




$$\begin{pmatrix} 0.0 \\ 1.0 \\ 2.0 \\ 2.0 \\ 0.0 \\ 1.0 \\ 1.0 \\ 0.0 \\ 2.0 \\ 1.0 \\ 1.0 \\ 2.0 \\ 2.0 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 1 & 1 \end{pmatrix} \begin{pmatrix} 0.0 \\ 1.0 \\ 1.0 \\ 0.0 \\ 0.0 \\ 1.0 \\ 1.0 \\ 1.0 \\ 1.0 \\ 0.0 \end{pmatrix}$$

Wish to solve inverse: $\vec{c} = \mathbf{G}_{wc}^{-1} \vec{w}$. However, $N_{cells} \approx N_{wires}^2$
 \Rightarrow as N_{cells} grows, more unknowns (\vec{c}) than knowns (\vec{w})!

Solution Attempt 2 - blob-level

Goal: **reduce matrix size** and **remove ambiguity** by making “blobs” of cells.



- 1 Select set of all cells with all three wires “hit”.
- 2 Partition set into spatially contiguous subsets: “**blobs**”.

Equation to solve is like the previous:

$$\vec{w}_b = \mathbf{G}_{wb} \vec{b}$$

$\vec{c} \rightarrow \vec{b}$ vector of charge in each blob.

$\mathbf{G}_{wc} \rightarrow \mathbf{G}_{wb}$ the wire-blob adjacency matrix for the slice.

$\vec{w} \rightarrow \vec{w}_b$ vector of charge on all wires associated with a blob.

But now **more favorable numerology**: $N_{blob} \lesssim N_{wb}$

Another wrinkle: charge uncertainty

Measure of the drifting charge by a wire has uncertainty:

- Environmental, electronic and thermal noise.
 - Can be correlated across wires/channels/chips/boards/etc.
- Statistical uncertainty due to digitization.
- Systematic uncertainties from detector response deconvolution.


Compare measured wire charge (\vec{w}_{meas}) with expected (\vec{w}_{exp}) and form a χ^2 with a covariance uncertainty matrix V .

$$\chi^2 = (\vec{w}_{meas} - \vec{w}_{exp})^T V^{-1} (\vec{w}_{meas} - \vec{w}_{exp})$$


“It can be shown” that minimizing this χ^2 is equivalent to inverting \mathbf{G}_{wc} (or \mathbf{G}_{wb}) matrix equation.

→ **This is a very CPU intensive, but critical step!**
(and, btw, we are still in the raw-data processing stage!)

The Payoff: imaged 3 GeV ν_e interaction



True energy depositions.



Wire Cell Imaging.

- **Excellent imaging** of major track features as well as isolated activity.
→ a static 2D view doesn't do it justice! [See it for yourself.](#)
- **Residual ambiguity** seen as wide blue patches.
→ **Inherent in LArTPC technology**, Wire Cell just makes it evident.
 - due to activity running parallel to the wire plane in one time slice.
 → Battling this will likely require (even more) **CPU-intensive algorithms** such as **iterative reconstruction**.

Pattern Recognition

Our current **post-imaging** approach:

- ➊ **cluster** together blobs contiguous in space and time (slice).
- ➋ **track** a line through a cluster.
- ➌ **categorize** success/failure of track to account for the cluster's charge distribution.

Some categories:

track cluster is well characterized by the track.

shower cluster consistent with an EM/hadronic shower.

short cluster appears to be a “short track” (eg, δ -ray).

undefined no well-suited categorization.

- This is an **active area of development** for Wire Cell.
- **Collaborate with machine learning experts?**

Physics-level Reconstruction

This development is just beginning, still largely conceptual.
The usual obvious goals:

- associate clusters with a particle trajectory
 - determine particle type
 - determine momentum, range, dE/dx and other kinematics.
- feed to Physics analysis.

Introduction

Technique

Software Design

Bee Display

Prototype

Toolkit

Future Plans

Wire Cell Software Ecosystem

Wire Cell breaks up into three main parts:

visualization the “Bee” web application (Chao Zhang)

prototype reconstruction algorithms, initial proof of principle
(Xin Qian)

toolkit “production” toolkit for LArTPC reconstruction (bv)

Bee: an interactive 3D visualization system

Bee features:


- Displays results of different reconstruction algorithms.
- Shows “true” particle trajectories from simulation.
- Very good for developers to debug and compare different methods and for users to gain Physics/LArTPC intuition.
- Implemented in **JavaScript/WebGL + Django**.
- Works on popular desktop and mobile browsers.
- Requires hardware WebGL support.
- JSON data file format, easy to implement schema.
- Supports drag-and-drop **user file uploads**.
- New features almost daily (**Chao Zhang!**)

Bee: Select and upload event sets

Instructions

.....
to upload
ore or clic
op the fi
.....


Interactive 3D visualization



Try it yourself! <http://www.phy.bnl.gov/wire-cell/bee/>

Wire Cell Working Prototype

- **Very successful proof of principle!**
- Currently **leads the state of the art** in LArTPC reconstruction techniques.
- Amazingly fast development (**Xin Qian!**)
 - Work started only ~5 months ago!



Colors indicate identified tracks and showers.

Compromises to Remedy

Some compromises were accepted to enable the **rapid prototyping** and now need “back-filling”. Need to remedy:

- No support for long-term, multi-developer contributions.
- Many monolithic, single-threaded applications.
- Hard-coded “magic” numbers → configuration.
- Rigid code structure and hard-coded workflows.
- External software integrated only through file exchange.
- Need comprehensive CPU/memory optimization campaign.
- Some high-level tests, need coverage/granularity.
- Intimate dependency on the large ROOT¹ toolkit.


Decision: rewrite rather than refactor →

¹<https://root.cern.ch/>

Wire Cell Toolkit - “Prototype 2.0”

High level overview:

- Includes **packaging and build** system.
- Comprehensive **API** via abstract base classes.
- Mindful of **dependencies**, external and internal.
- Built-in wire and cell **geometry descriptions** or load from file.
- Includes simple **LArTPC detector simulation**.
- Rewrite implementations of core **prototype algorithms**.
- Abstracted **execution model**, step towards parallelism.
- Support for native and foreign **file I/O**.
- Follow “**toolkit**” software paradigm.



Good progress along all fronts but still much work to do.

Source Repositories and Build

- Code in GitHub [WireCell](#) organization.
- Code aggregation with `git submodule`.
- A simple, customized `waf`-based build.²

```
$ git clone git@github.com:WireCell/wire-cell.git
$ cd wire-cell/
$ git submodule init
$ git submodule update
$ ./wcb --prefix=/path/to/install configure build install
```

Builds, tests and installs:

- toolkit shared libraries + header files,
- no main applications yet but many, unit+integration tests

Also: [Doxygen code reference](#) and [MkDocs user documentation](#).

²`wcb` = `waf` + extra Waf tools for Boost, ROOT, Eigen and Wire Cell packaging


Wire Cell Interfaces

The `wire-cell-iface` package:

- API composed of abstract base classes covering:
 - **Data model** (wires, cells, frames, slices, tracks, ...)
 - **Active components** (blob maker, matrix solver, clustering, ...)
- Pervasive use of C++ `shared_ptr<>` for (mostly) **worry-free memory management**.
- Dynamic instance lookup via `NamedFactory` pattern allows for a **plugin architecture**.
- Initial support for **user configuration** system based on Boost property trees.
- Initial support exists for an **abstract execution model**.
 - more on this coming up.

Wire Cell Simulation

- Provided by the `wire-cell-gen` package.
- Simple implementation of the 4 “D”s of LArTPC: **deposit, drift, diffuse, digitize**.
- Granular, use as little or as much as needed.
 - Parts can be replaced by **external simulation**, or bypassed entirely with **real DAQ data**.
 - **External integration** via file I/O or API calls.
- Provides reference implementation to guide integrator/developers of external frameworks.
- Useful for quickly generating data to feed unit tests.



- Design directly influenced by **CSC Seminar on VisTrails, March 2013!**
- Data flows through a graph made from:
 - vertices:** computational units / algorithms
 - edges:** data queues of a given type
- Streamed processing minimizes RAM usage.
- Graph-level “programming” to define workflows.
- Thread-safe queues \Rightarrow parallel processing.
- Graph execution machinery swappable: uniproc, multiproc, or distributed (MPI) parallel.
- Encourages isolated, targeted development and testing of each compute vertex.
- Feedback loops to implement iterative workflows.
- Instrument graph to collect performance data.




Abstract execution model

Want ability to change execution model while leaving “real” algorithm code untouched!

Each graph vertex made from **three layers**.

From outside to in:

- 1 Execution model layer implementing the data flow control
 - sync'ed load-and-flush, async, multi-processing, message-passing, ...
 - some implementations identified
- 2 Adapter wraps actual algorithm and presents uniform interface to execution model layer.
- 3 Algorithm implementation and interface left unrestricted except “no shared globals” (thread safety).



Introduction

Technique

Software Design

Future Plans

- Algorithm Development

- Bee 2.0

- Parallel Wire Cell

Future Plans

...in which I frequently beg for help :)

Pattern Recognition

Pattern recognition stage is turning out to be “hard”.

- Many challenges are **inherent to the LArTPC technology**.
 - Wire Cell has just allowed us to reach them.
 - Current implementation is largely **heuristic based**
 - many “cross-algorithm assumptions”
 - **large solution space** with **many corner cases**.
 - human Physicists find good but **per-event/ad-hoc solutions**.
 - how to teach them to the computer?
- Our attempt: Bee 2.0 (next slide)

We have a few ideas but we feel we may be entering to the area of “real” computer science problems like **machine learning** and maybe **machine vision**.

And we are not experts.

Help Wanted!

Bee 2.0: Human-directed Automated Reconstruction

Interim Solution: inject human pattern recognition.

- 1 Bee will let humans **interactively inject commands**, eg:
 - 1 "Join cluster X to cluster Y"
 - 2 "Remove blob AAA from cluster Z"
 - 3 "Now, rerun automatic reconstruction"
- 2 Package human intentions into command objects.
- 3 Send to back-end, interpreting and execution inside a Wire Cell service.
- 4 Send results back to user, possibly iterate.
- 5 Record everything for later replay and to mine for potential Wire Cell improvements.

Think: "**Physics-PhotoShop**" + "**Google Analytics**".

We have a rough conceptual architectural design but it will be a **rather big** system. \Rightarrow **Help wanted!**

Wire Cell Production Processing

A **1-2 punch** for today's computing:

- Need **high throughput** to keep up with detector data rates.
- Need **high performance** due to CPU-heavy algorithms.

Data volumes:

MicroBooNE ~630 TB/year (compressed)

protoDUNE $\mathcal{O}(1\text{PB})$ total, (ZS, uncomp.)

DUNE 40kt $\mathcal{O}(\text{PB})/\text{year}$ (ZS, uncomp.)

- Okay, it's not quite ATLAS-levels
 - Run 1 produced ~3PB/year raw.
- But: **Wire Cell needs all raw wire waveforms** as input
 - No "Level-X" triggers to save us.
 - Waveform thresholding, but otherwise everything.
- Current speed is **30 minutes – 1 day / event!**
 - **Major caveats:** we expect some optimization still to be had, we will parallelize the code, etc, etc,
 - But: ... ***GULP*!**

(with apologies to Frank
Wuerthwein)



Taking Wire Cell Parallel

Our plans for **parallel processing** mostly involve implementing and testing different **parallel execution models**:

- 1 **Boost.Pipeline** nice and simple threaded data flow package. Easy to start with but not yet officially part of BOOST (longevity concerns?).
- 2 **Intel TBB** a common and well supported library with a fantastic looking “flow graph designer” tool. Some effort needed to learn/adopt.
- 3 **MPI/ZeroMQ** or other ways to run Wire Cell as multiple-node distributed process (maybe on **HPC**). Big learning curve for us.
- 4 Finally we will investigate potential for **GPU/Phi** hardware acceleration of bottlenecks. Even bigger learning curve!

This level of parallelism is **new ground for us**³.

Help wanted!

³This is a subject of a pending BNL/Computing LDRD proposal.

Summary

- The **Wire Cell working prototype LArTPC reconstruction** method and software has been developed.
 - already producing some of the **world's best results**,
 - technical and performance improvements still needed.
- The **Bee interactive 3D event visualization** application developed, critical for understanding LArTPC and developing Wire Cell. Plans to evolve to “human-directed automated reconstruction” system.
- The **Wire Cell Toolkit** for LArTPC reconstruction partly complete and will be the basis for long term development including investigations into parallel processing.
- To be successful we must wade into **new waters** (for us) of parallel processing, hardware acceleration, computing science and mathematics.

→ **Expert input, help and collaboration are most welcome!**

Wire Cell on the web

- home page:
<http://www.phy.bnl.gov/wire-cell/>
- Bee entry page:
<http://www.phy.bnl.gov/wire-cell/bee/>
- prototype user manual:
<http://bnlif.github.io/wire-cell-docs/>
- prototype repositories:
<https://github.com/BNLIF>
- toolkit user manual:
<http://wirecell.github.io/wire-cell-docs/>
- toolkit code reference:
<http://www.phy.bnl.gov/wire-cell/doxy/html/>
- toolkit repositories:
<https://github.com/WireCell>